



BITS Pilani

Cloud Computing

Session 4

Containers

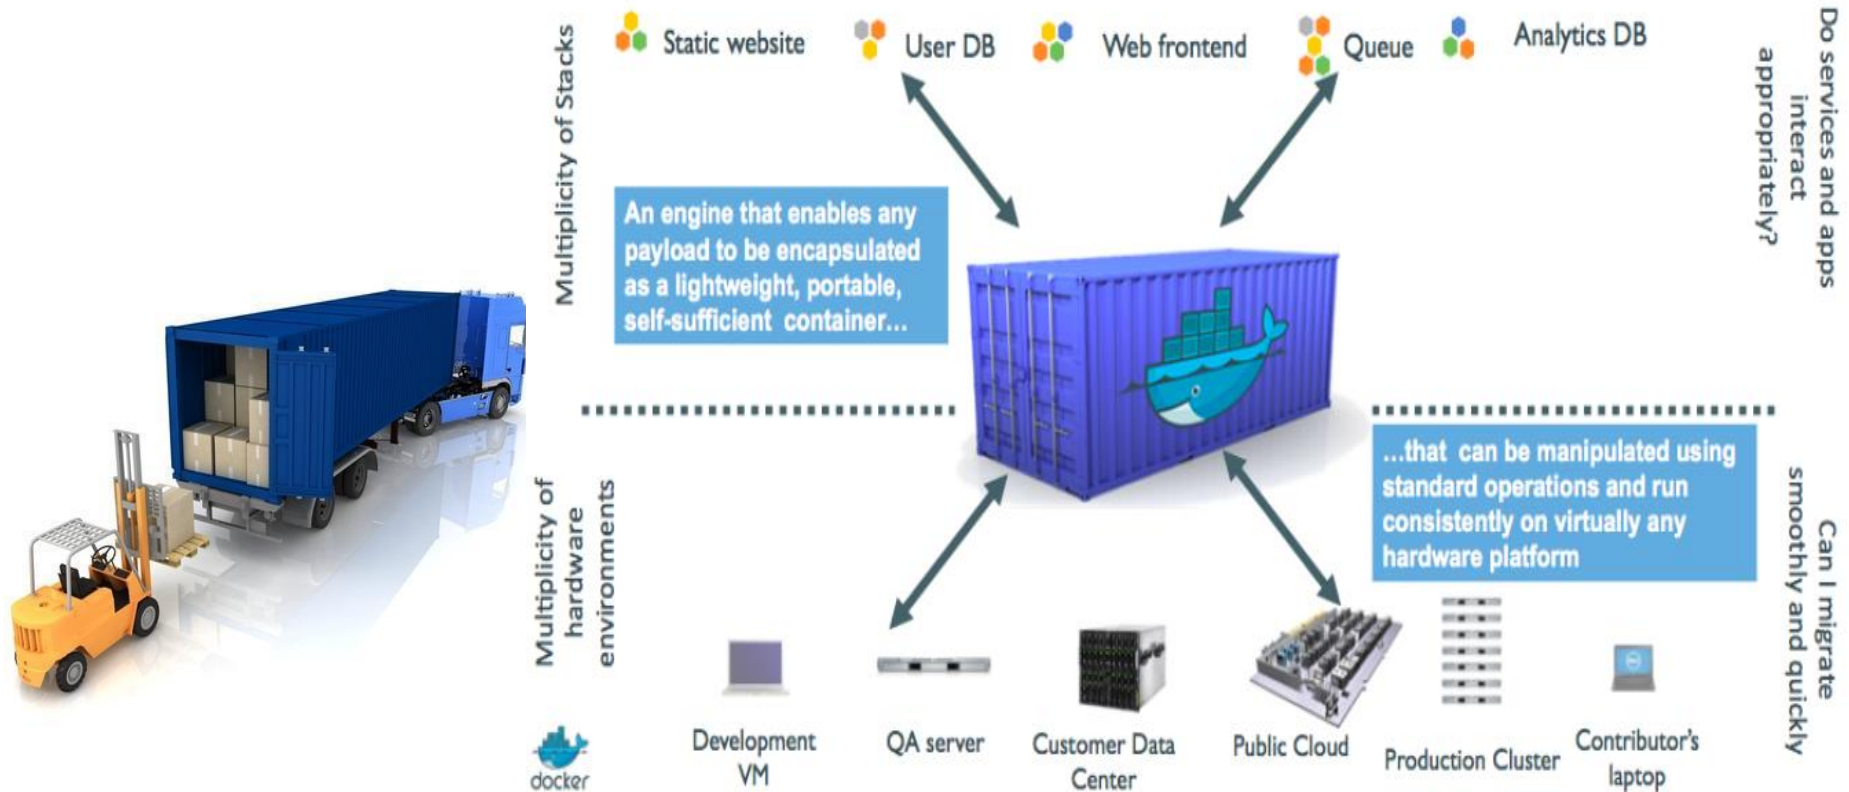
Shwetha Vittal
shwetha.vittal@pilani.bits-pilani.ac.in

Agenda

- ❑ What are containers?
- ❑ Namespaces
- ❑ Cgroups
- ❑ Virtual Machine vs Containers
- ❑ Docker

What are Containers?

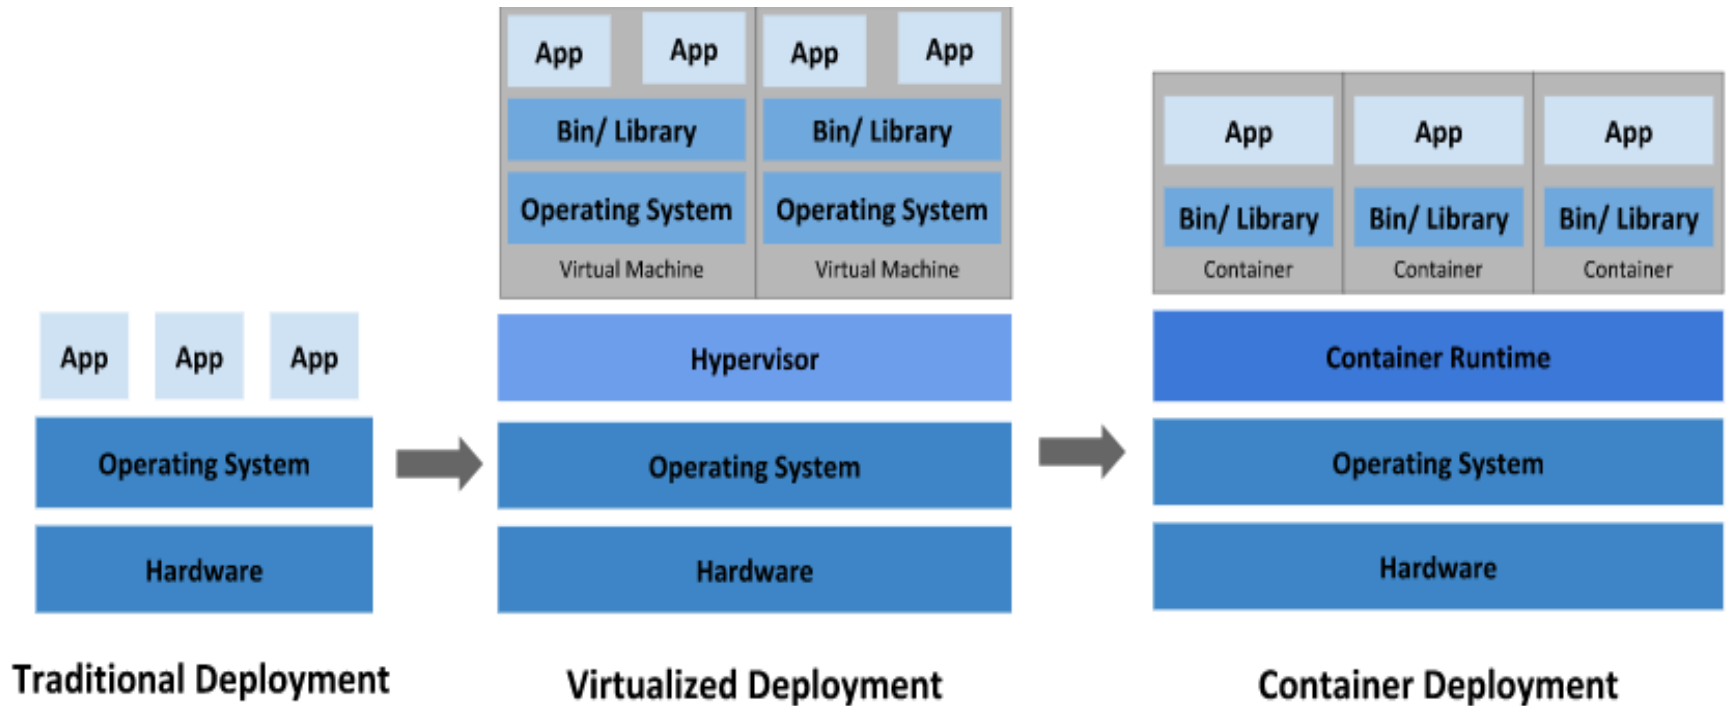
A shipping container system for applications



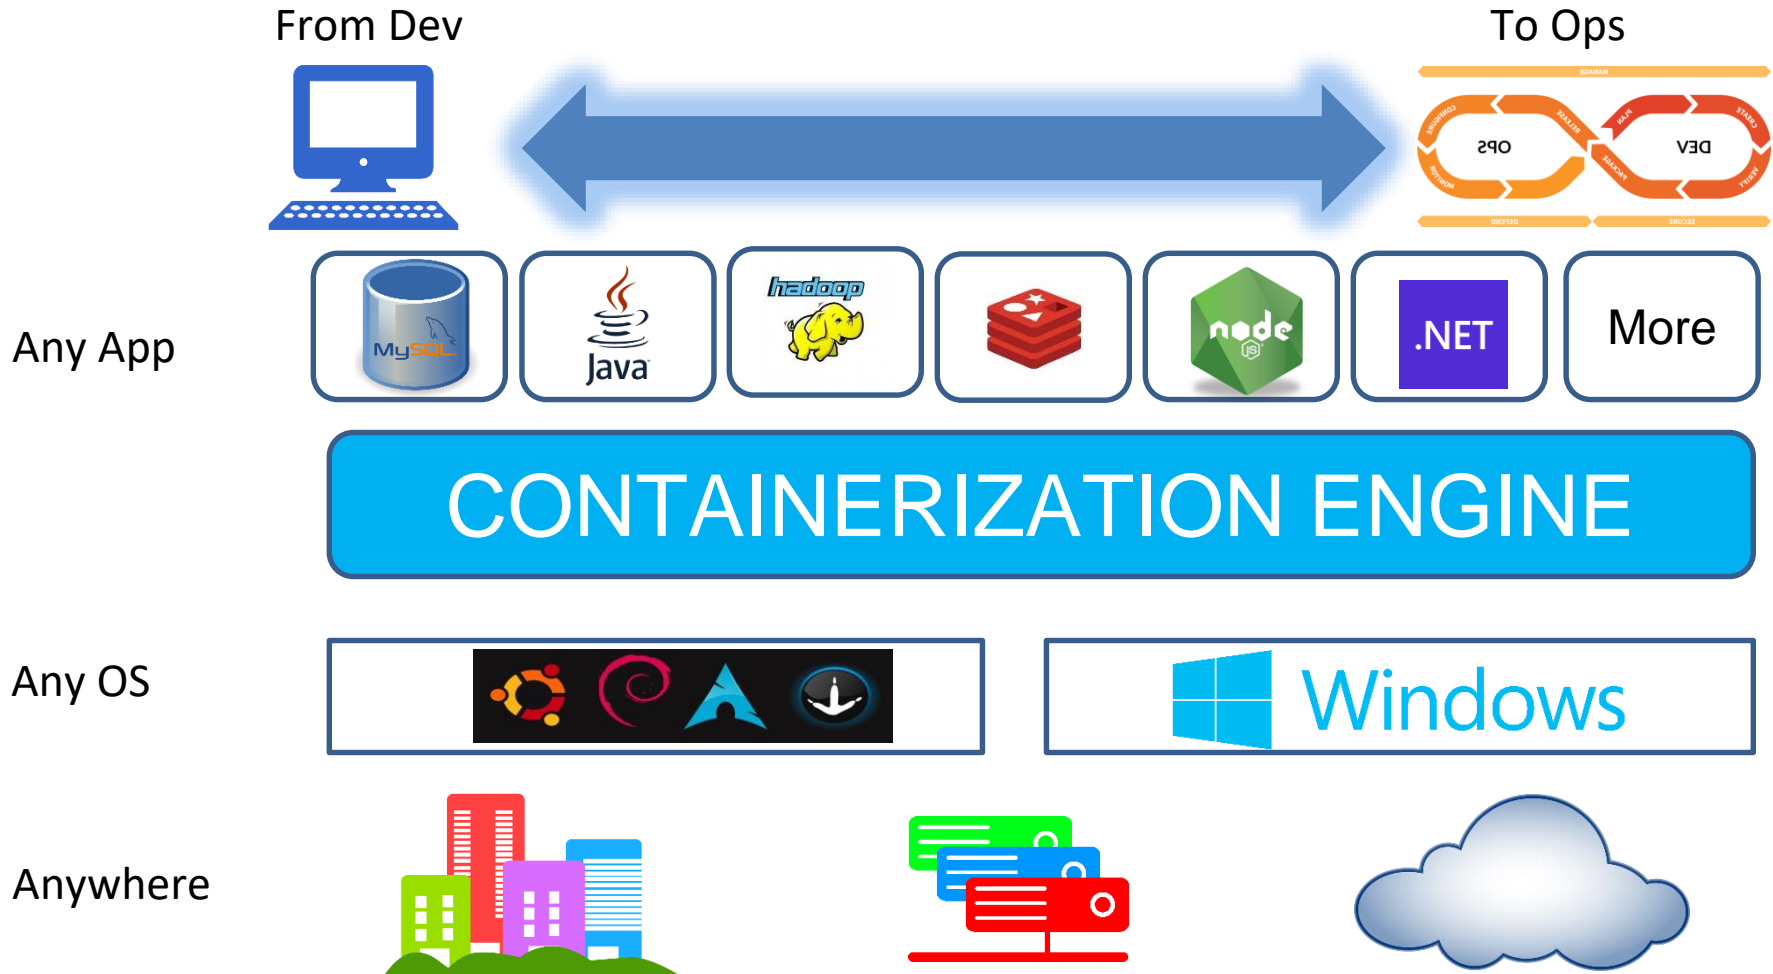
What are Containers ?

- Lightweight virtualization mechanism
- A software container is a standardized package of software.
- Everything needed for the software to run is inside the container
- The software code, runtime, system tools, system libraries, and settings are all inside a single container
- Managed by the OS kernel running on the host system
- Has its own isolated memory, CPU, storage, process table, and networking interfaces
- Faster provisioning for newer applications

Going back in Time to Now




What are Containers?



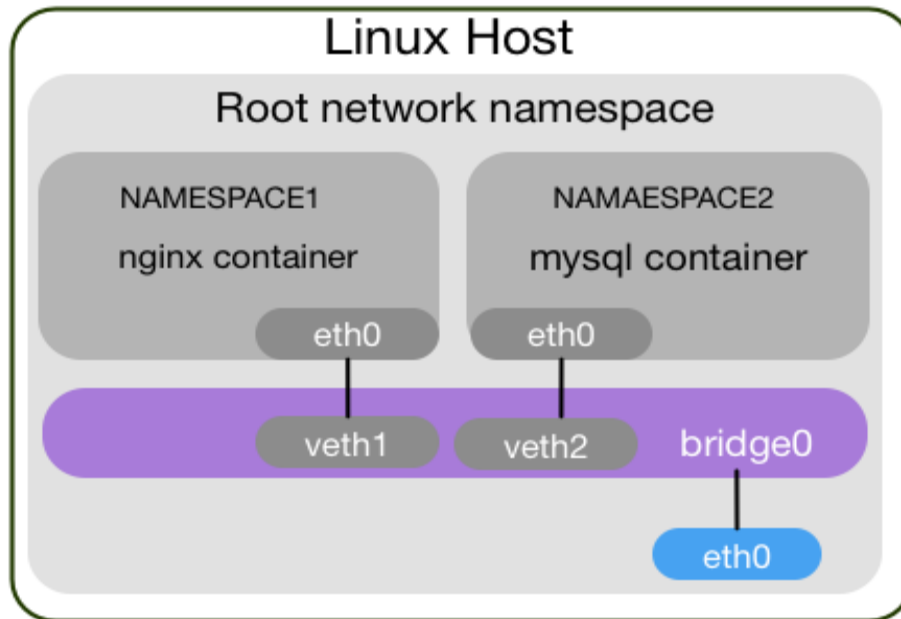
Containers

Containers are powered by two underlying Linux Kernel technologies

- Namespaces
- Cgroups



Namespaces



- Kernel mechanism for limiting the visibility that a group of processes has of the rest of a system
- Namespace merged into Linux 3.8
- Limit Visibility
 - Process trees - PIDs
 - Network interfaces
 - User IDs
 - Filesystem mounts

Types of Namespaces

- Wrap a particular global system resource in an abstraction
- **Illusion:** Makes it appear to the processes within the namespace that they have **their own isolated instance of the global resource**.
- 6 Main Namespaces
 - Mount namespace
 - UTS namespace
 - IPC namespace
 - PID namespace
 - Network namespace
 - User namespace

```
$ ls -l /proc/13/ns
total 0
lrwxrwxrwx 1 root  root      0 Feb  6 09:57 cgroup -> cgroup:[4026531835]
lrwxrwxrwx 1 root  root      0 Feb  6 09:57 ipc  -> ipc:[4026547635]
lrwxrwxrwx 1 root  root      0 Feb  6 09:57 mnt  -> mnt:[4026547631]
lrwxrwxrwx 1 root  root      0 Feb  6 09:57 net  -> net:[4026547638]
lrwxrwxrwx 1 root  root      0 Feb  6 09:57 pid  -> pid:[4026547636]
lrwxrwxrwx 1 root  root      0 Feb  6 09:57 user -> user:[4026531837]
lrwxrwxrwx 1 root  root      0 Feb  6 09:57 uts  -> uts:[4026547632]
```

Mount Namespace

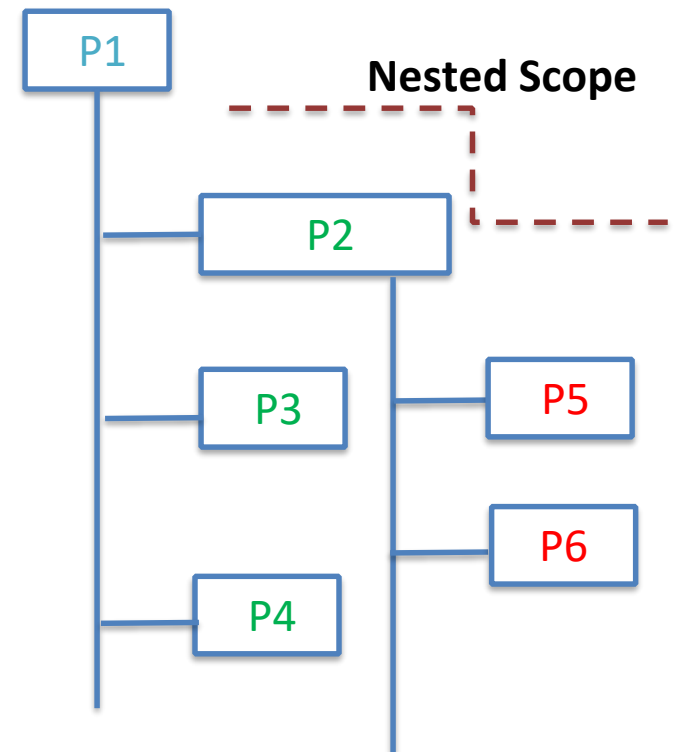
- Isolate the set of filesystem mount points seen by a group of processes.
- Processes across different mount namespaces (ns) have different views of the filesystem hierarchy.
- `mount()`, `umount()`
- Shared or Private mount points
 - Shared mount points propagated to all namespaces across process hierarchy / other processes.
 - Private is not
- Every container – has a custom root file system to start with.
- Any new child process without any shared ns by its parent, will start with empty root filesystem

IPC Namespace

- Isolate certain inter process communication (IPC) resources
 - System V IPC objects
 - POSIX message queues.
- Have a private set of IPC objects (sem, shm, msg) inside namespace.

PID Namespace

- Isolate the process ID number space.
- Processes in different PID namespaces can have the same PID.
 - Containers can be migrated between hosts while keeping the same process IDs for the processes inside the container.
 - Allow each container to have its own init – PID 1
- Nested Scope
 - Ancestor->...->Parent -> child



Network Namespace

- Provide isolation of the network resources
- Each network namespace has its own network devices, IP addresses, IP routing tables, /proc/net directory

List the network ns(es)

```
ubuntu@ip-172-31-31-148:~$ ls -l /var/run/netns;  
total 0  
-r--r--r-- 1 root root 0 Feb  6 10:19 mynetworkns  
-r--r--r-- 1 root root 0 Feb  6 10:17 testns
```

```
ubuntu@ip-172-31-31-148:~$ ip netns  
testns  
mynetworkns
```

Adding network ns

```
ubuntu@ip-172-31-31-148:~$ sudo ip netns add mynetworkns  
ubuntu@ip-172-31-31-148:~$
```


Network Namespace

Network interfaces on host

```
ubuntu@ip-172-31-31-148:~$ ip link list
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: enx0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 0a:ff:ed:60:2c:e5 brd ff:ff:ff:ff:ff:ff
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT group default
    link/ether 02:42:75:31:cf:1a brd ff:ff:ff:ff:ff:ff
```

Network interfaces inside ns

```
ubuntu@ip-172-31-31-148:~$ sudo ip netns exec mynetworkns ip link list
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```



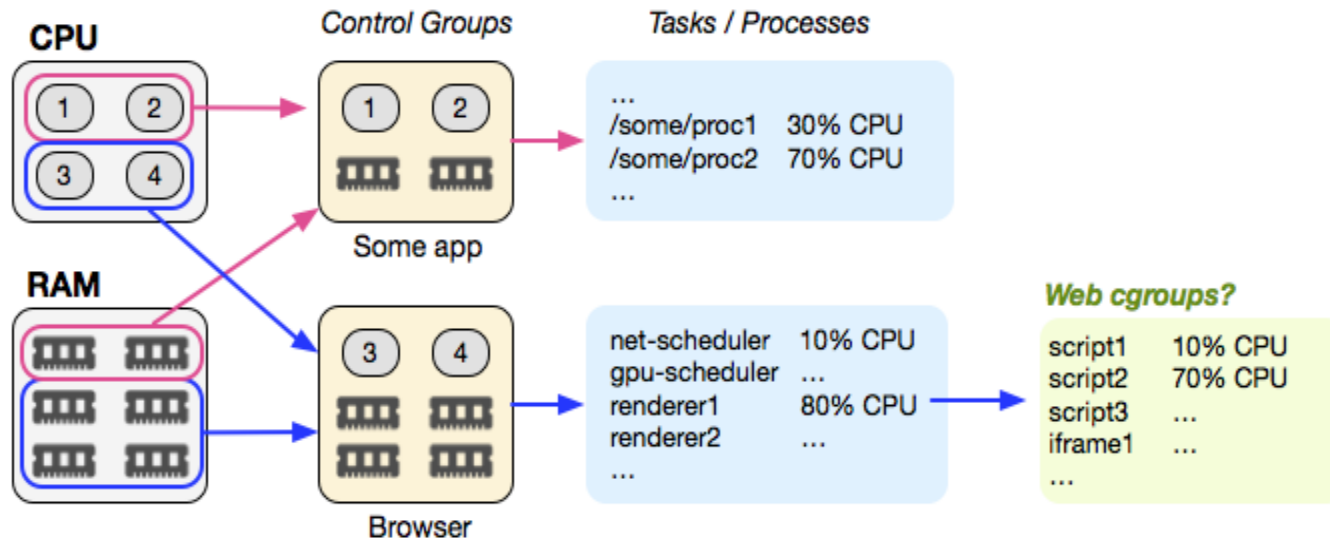
User Namespace

- Isolate the user and group ID number spaces
- A process's user and group IDs can be different inside and outside a user namespace.
 - A process has full root privileges for operations inside the user namespace,
 - But is unprivileged for operations outside the namespace.

Namespace APIs

- System Calls
 - `clone()`: Create a new process and place it into a new namespace.
 - `unshare()`: Creates a new namespace and places calling process into it.
 - `setns()`: Join an existing namespace.
- Commands
 - `lsns` - all namespaces in the system
 - `/proc/PID/ns` - which namespace a process belongs to.

Cgroups



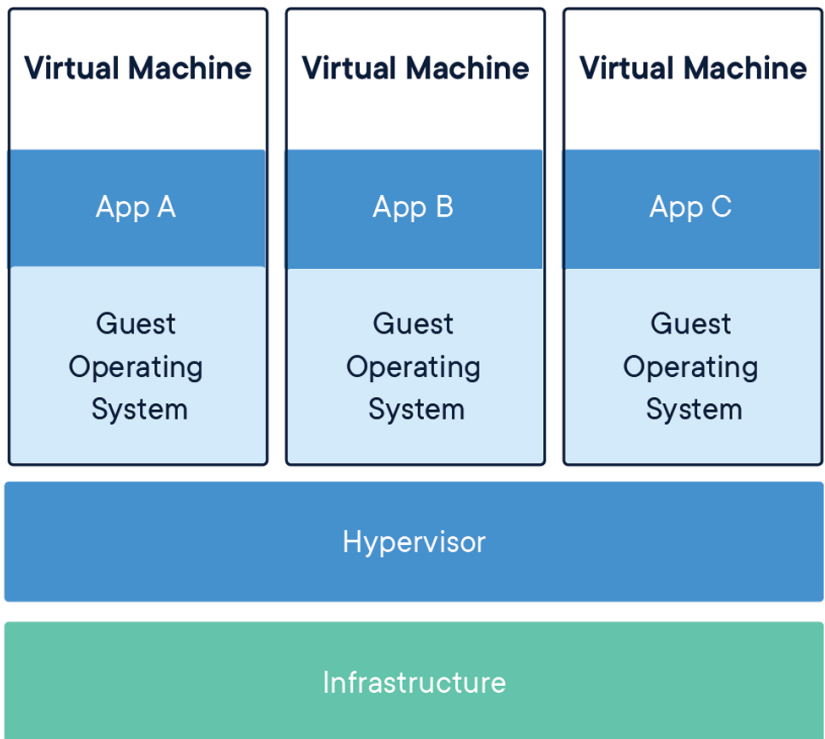
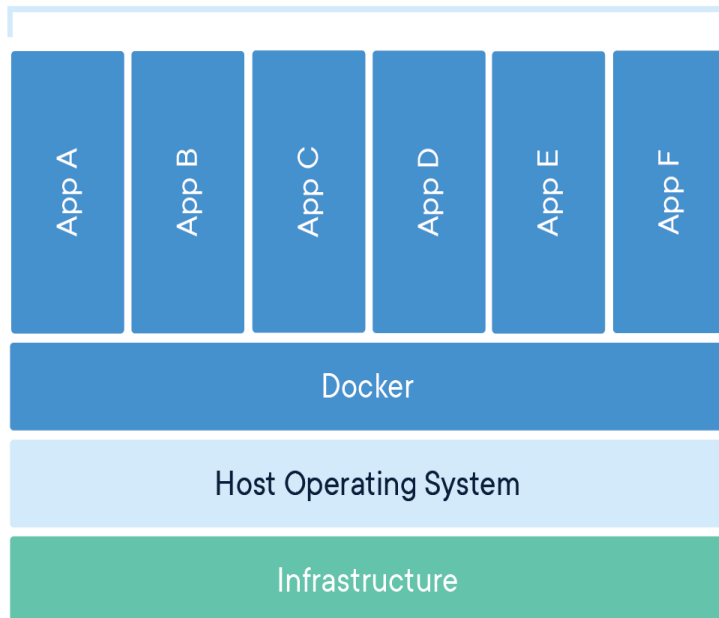
cgroups – Control groups

- A kernel mechanism for limiting and measuring the total resources used by a group of processes running on a system
- Processes can be applied with CPU, memory, network or IO quotas

Cgroup merged into Linux 2.6.24

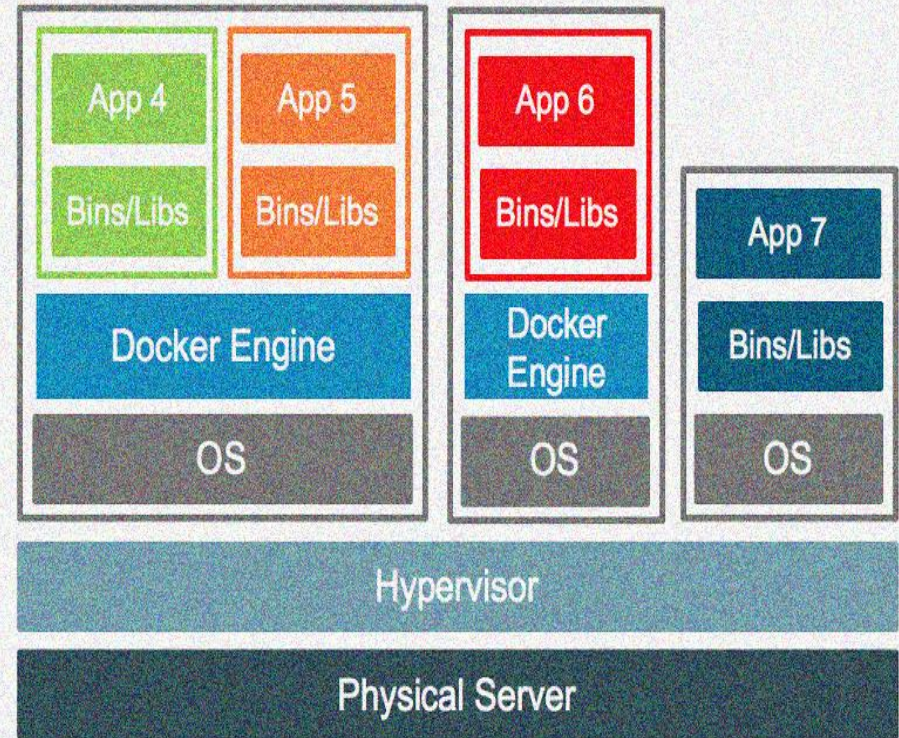
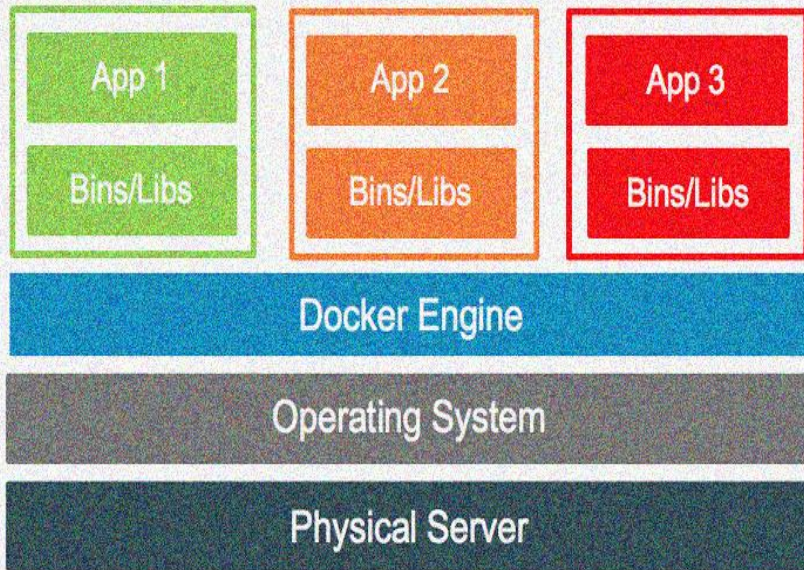
Containers Vs Virtual Machines

Containerized Applications



Containers on Virtual Machines ?

Your Datacenter or VPC

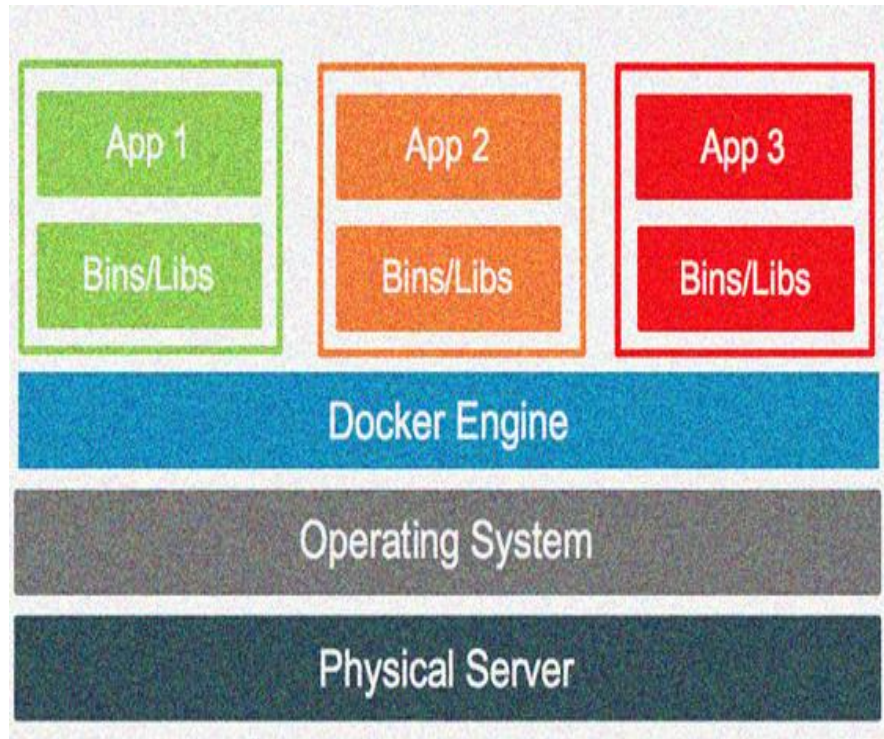




Docker

Docker Platform

- Docker is an open platform
- Docker separates applications from hardware infrastructure
- **Containers** are used to package and run an application

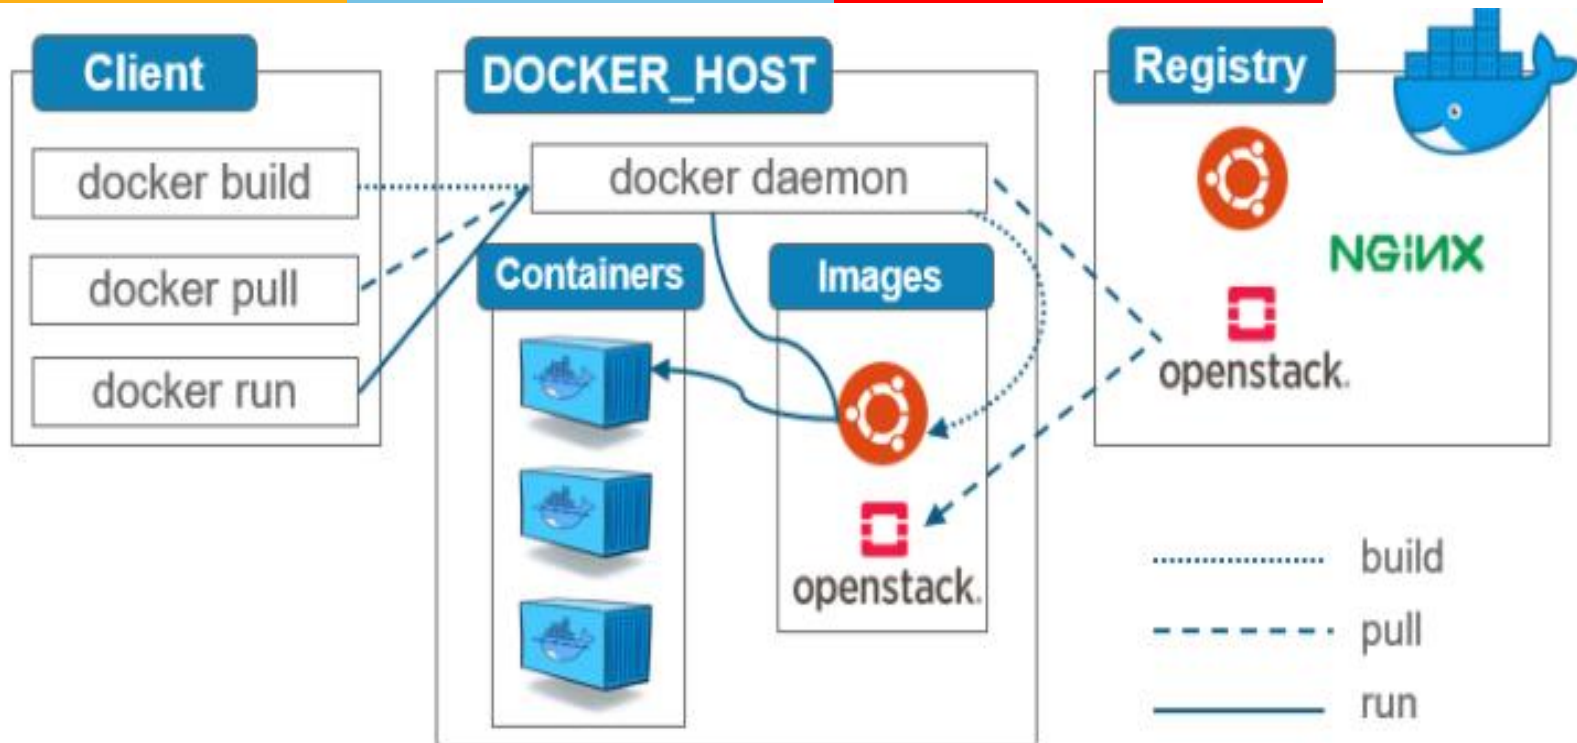


- A single host can run many containers simultaneously
- Containers are lightweight and contain everything needed to run the application

Docker Platform

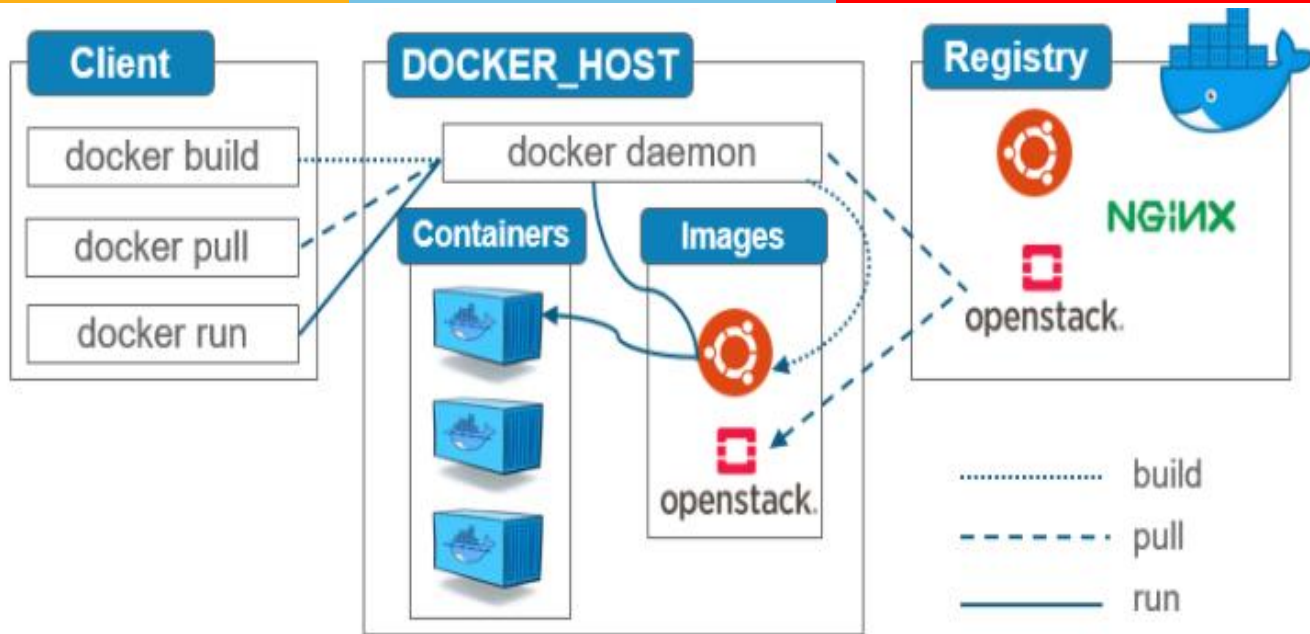
- Docker provides tooling and a platform to manage the lifecycle of your containers:
 - Develop application(s)
 - Distribute & test
 - Deploy into production environment, as a container or an orchestrated service.
- Containers are great for continuous integration and continuous delivery (CI/CD) workflows.

Docker Architecture



- Docker uses a client-server architecture.
- The Docker daemon
- The Docker client
- The Docker client and daemon communicate using a REST API, over UNIX sockets or a network interface.

Docker Architecture



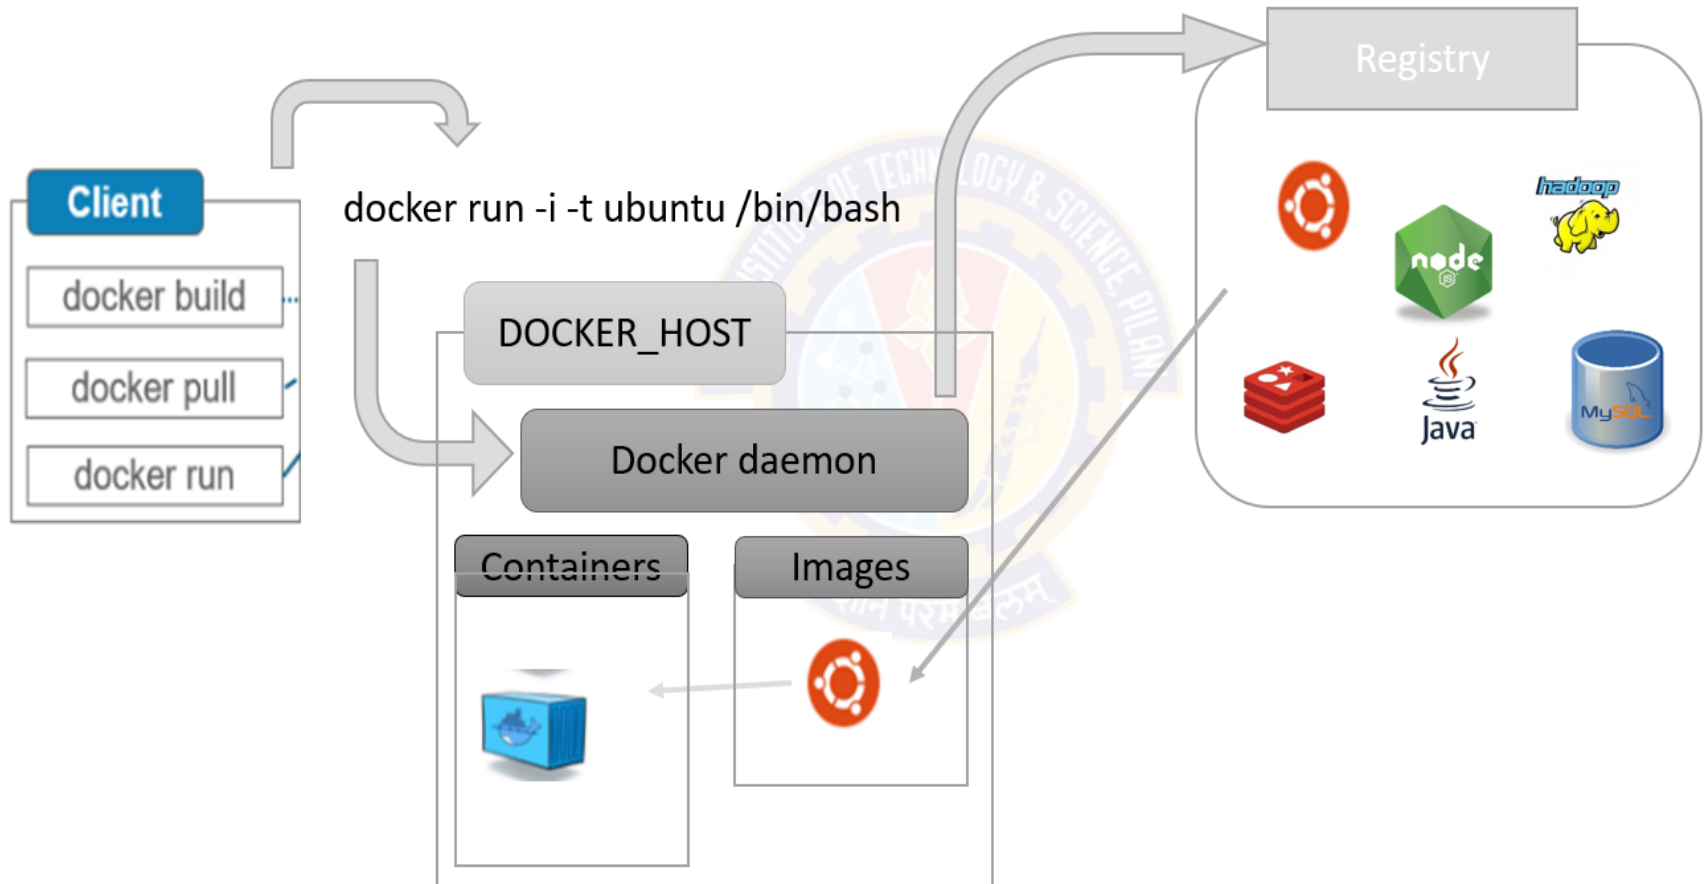
- **Docker Registries**

- Stores Docker images.
- Docker Hub is a public registry that anyone can use.
- Docker is configured to look for images on Docker Hub by default

- **Docker Objects**

- **IMAGES:** An *image* is a read-only template with instructions for creating a Docker container.
- **CONTAINERS:** A container is a runnable instance of an image

Running a Docker Container



Docker Commands

- A container is a runtime instance of a docker image
- Create and run a container from an image, with a custom name:

`docker run --name <container name> <image name>`

docker run --name mylinuxserver Ubuntu

- Run a container with and publish a container's port(s) to the host.

`docker run -p <host port>:<container port> <image name>`

docker run -p 8080:80 nginx

- Run a container in the background

`docker run -d <image name>`

docker run -d -p 8080:80 nginx

Docker Commands

- Start or stop an existing container:
docker start/stop <container name> (or <container id>)
docker stop 11ed (or mynginx)
- Remove a stopped container:
docker rm <container name> (or <container id>)
docker rm -f 11ed (or mynginx)
- Open a shell inside a running container:
docker exec -it <container name> sh
docker exec -it myubuntu bash

Docker Commands

- Fetch and follow the logs of a container: ***docker logs -f <container name>***
- To inspect a running container: ***docker inspect <container id> (or) <container name>***
- To list currently running containers: ***docker ps***
- List all docker containers (running and stopped): ***docker ps --all***
- View resource usage stats: ***docker container stats***

Docker Image

- A lightweight, standalone, executable package of software
- includes
 - code,
 - runtime,
 - system tools,
 - system libraries
- Build an Image from a Dockerfile: ***docker build -t <image name>***
- List local images: ***docker images ls***
- Delete an Image: ***docker rmi <image name>***
- Remove all unused images: ***docker image prune***

Build & Run Customized Image

Dockerfile: File with instructions to build a docker container image

```
FROM ubuntu:latest
RUN mkdir /app
RUN apt update
RUN apt install vim g++ -y
WORKDIR /app
ENTRYPOINT ["/bin/bash"]
```

Hands On

```
docker build -t myappimage .
docker run -it myappimage
```

Summary – So Far

- Introduction to Containers - Lightweight Virtualization
- Key Building blocks
 - Namespaces
 - Provide isolation
 - 6 key namespaces: Mount, UTS, IPC, PID, Network, User
 - Cgroups
 - Resource limiting, Prioritization
 - CPU, memory
- Containers Vs Virtual Machines
- Docker Containers
 - Docker Architecture: Client-server, Objects, Registry
 - Building Customized image – Dockerfile

Lab - Web Service on Docker

1. Deploy nginx web service on docker on your PC or via

<https://labs.play-with-docker.com/>

1. Once it is executing, try accessing it using curl <ip>:port from terminal or if you have browser locally use localhost:port

Capture all your observations and discuss with peers and instructor.

References

- Namespaces in operation, part 1: namespaces overview [LWN.net]
- [Chapter 1. Introduction to Linux Containers | Red Hat Product Documentation](#)
- [Docker Engine | Docker Docs](#)
- [The Resilience of Virtual Machines: Why VMs Are Still Vital in the Age of Docker Dominance | Effective Programmer | Effective Programmer](#)