



**BITS Pilani**

# Cloud Computing

## Session 5

# Container Orchestration

Shwetha Vittal

# Agenda


---

- ❑ Need for Container Orchestration
- ❑ Kubernetes (K8s) Features
- ❑ K8s Objects
- ❑ Use cases of K8s
- ❑ What K8s is not?

# Why Container Orchestration

## Container

- Logically distinct piece of software
- Built, deployed, maintained, managed and, scaled on own without unduly affecting other parts of the system.



## Container Orchestration

- Bigger system requires multiple containers
- Need interaction among themselves.
- E.g application servers need to interact with backend DB.
- Come together as a single scalable, reliable, and resilient system
- **Container orchestration** comes in for rescue
  - E.g. Kubernetes

# Kubernetes – What is it ?

- Kubernetes (K8s) - An open source system for
  - Automating deployment: Dynamically pushing configuration files to running jobs;
  - Service discovery and load balancing
  - Auto-scaling
  - Management of containerized applications.




Image Courtesy: [Kubernetes](#)

[15 years of experience of running production workloads at Google](#)

Groups containers that make up an application into logical units for easy management and discovery.

# Kubernetes – What is it ?

---




Image Courtesy: [Kubernetes](#)

- High Availability service:
  - K8s provides you with a framework to run distributed systems resiliently.
  - It takes care of scaling and failover for your application

# Features of K8s

- **Automatic bin packing**
  - K8s can fit containers onto the specified nodes to make the best use of your resources.
- **Self-healing**
  - K8s restarts containers that fail, replaces containers, kills containers that don't respond to user-defined health check




Image Courtesy: Bytebytego

# Features of K8s

- **Service Discovery and Load balancing**
  - **Discovery:** Kubernetes can expose a container using the DNS name or using their own IP address.
  - **Load balancing:** If traffic to a container is high, K8s distributes the network traffic




Image Courtesy: Bytebytego

# Features of K8s



Image Courtesy: Bytebytego

## Automated rollouts and rollbacks

- Change the actual state of the deployed container to the desired state at a controlled rate.
- E.g.: Remove existing containers and adopt all their resources to the new container.



# Features of K8s




Image Courtesy: Bytebytego

- **Horizontal scaling**
  - Scale your application up and down with a simple command, with a UI, or automatically based on CPU usage.
- **Designed for extensibility**
  - Add features to cluster without changing upstream source code.

# K8s Nodes & Cluster

---

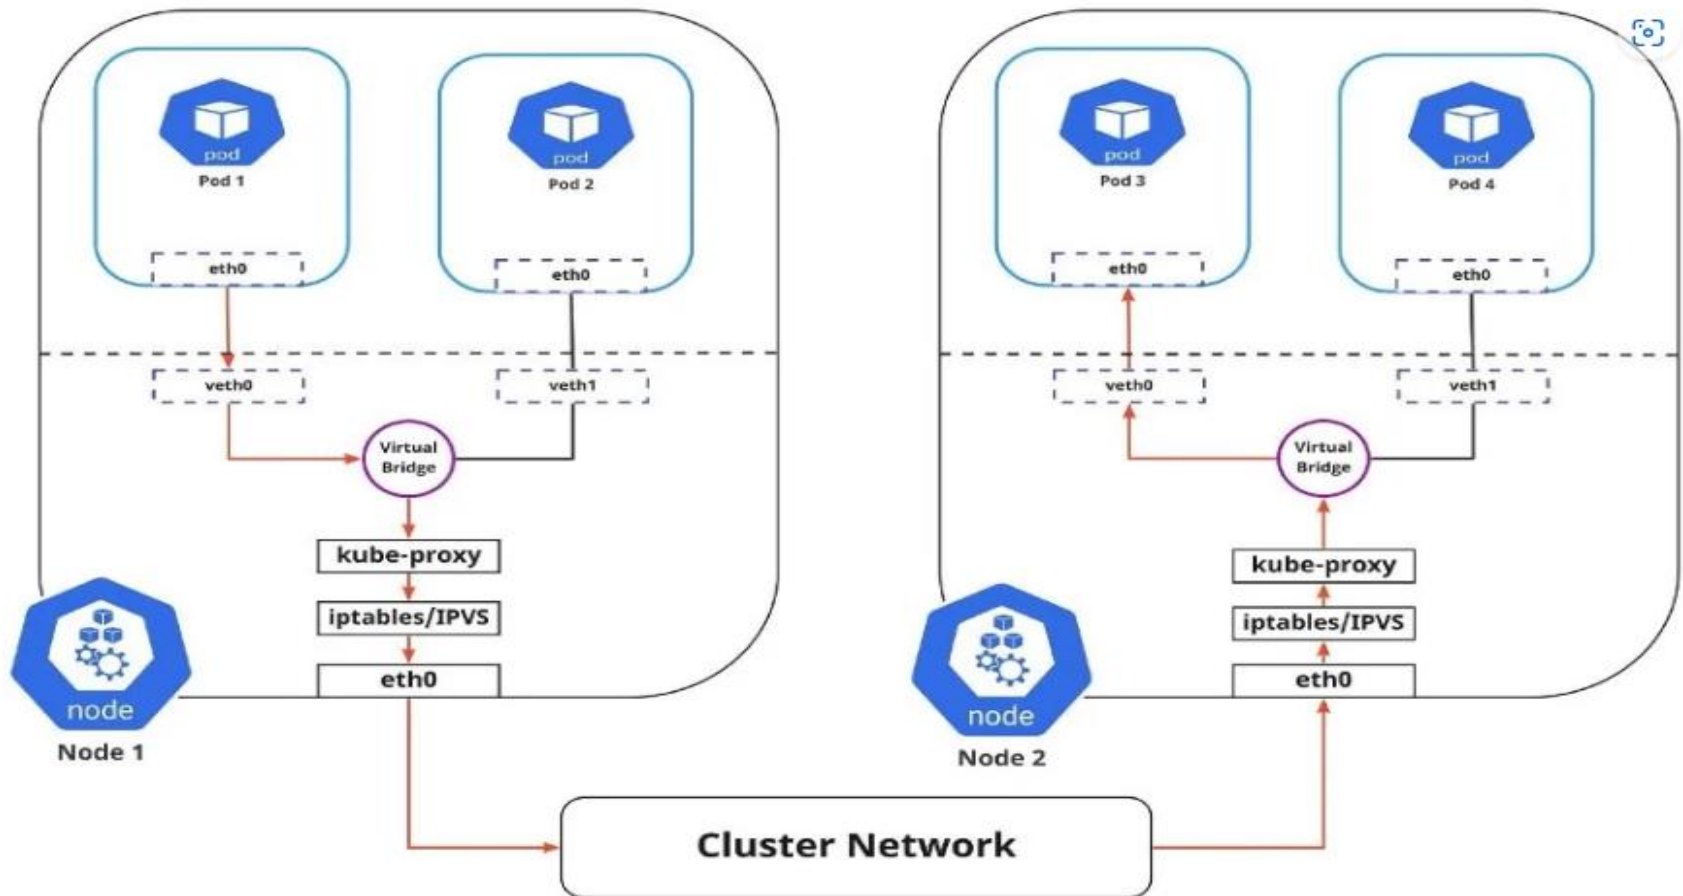
- **Node**

- A node is a host system, whether physical or virtual, with required components.
- Allows an execution of one or more containers inside it.
- Two types of Nodes
  1. Master
  2. Worker

- **Cluster**

- A K8s cluster consists of one or more nodes.
- K8s system components, that is connected to a network that allows it to reach other nodes in the cluster.

# K8s Nodes and Cluster



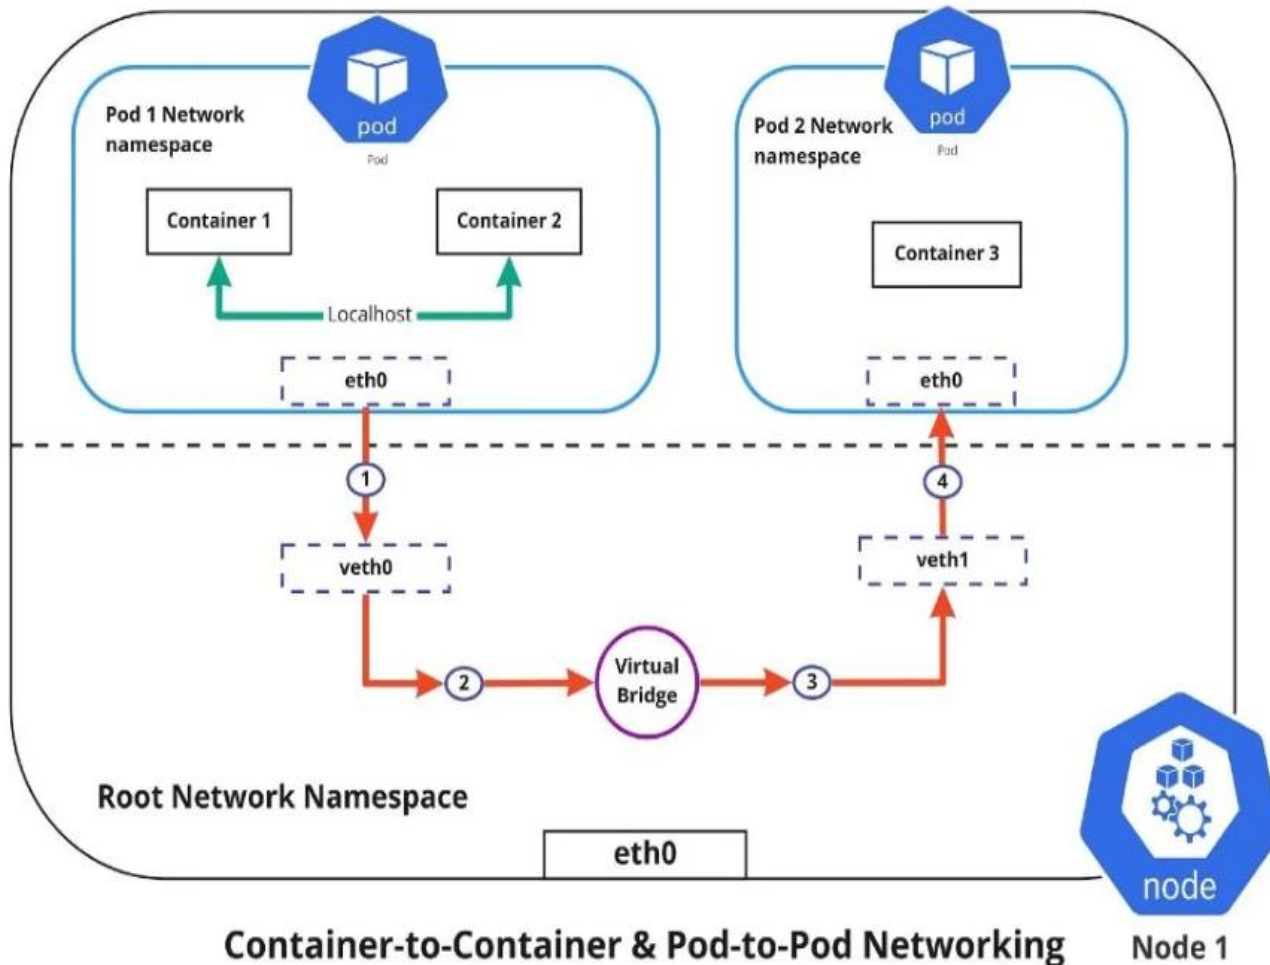
(Nived Velayudhan, CC BY-SA 4.0)

# K8s Network Model

---

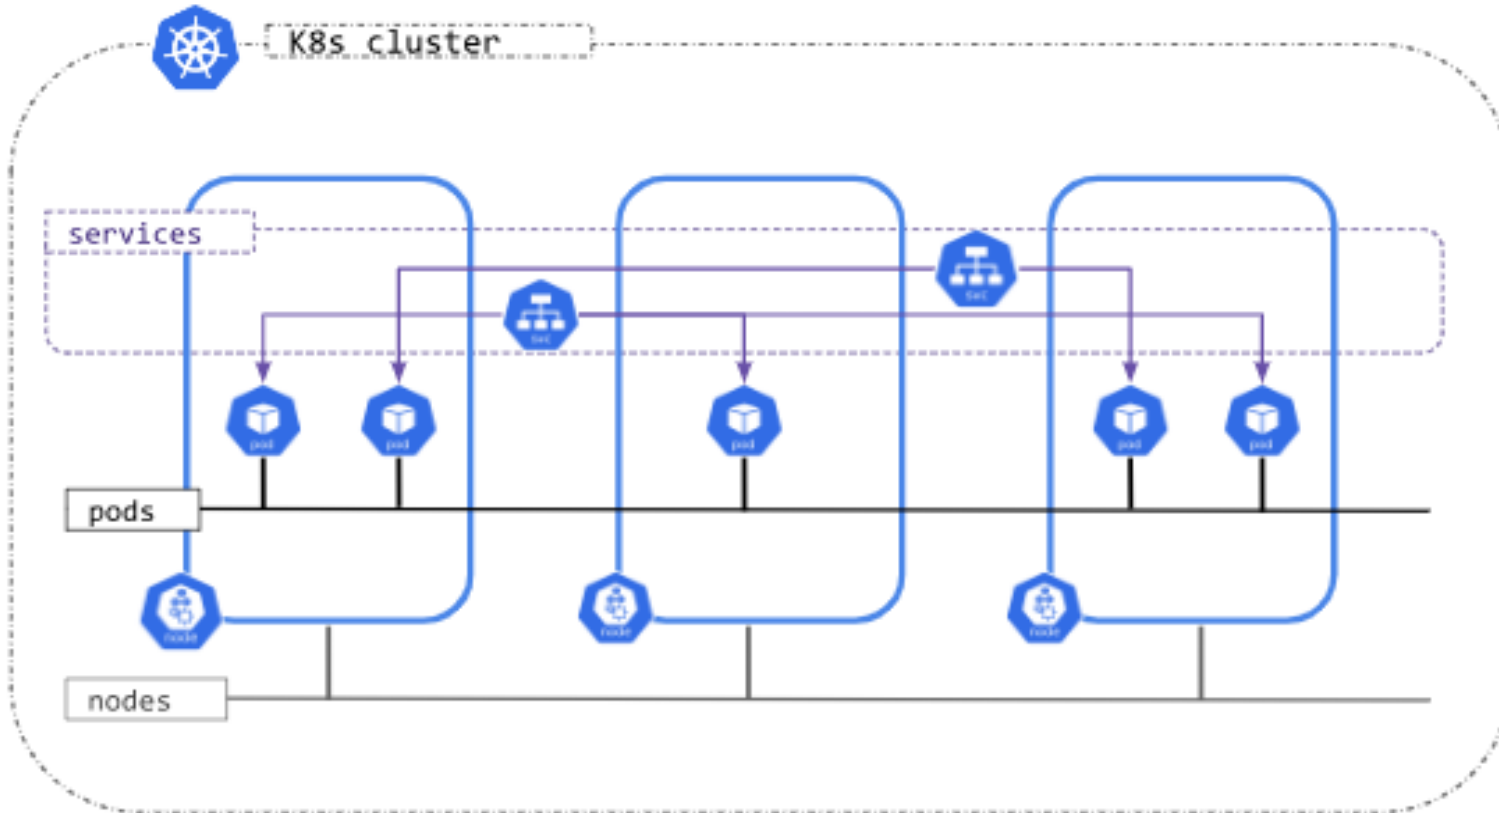
- **Pod**
  - Smallest deployable unit for executing back end code /application.
  - Gets its own unique cluster-wide IP address.
  - Has its own private network namespace which is shared by all of the containers within the pod.
  - The pod network handles communication between pods.

# K8s Network Model



(Nived Velayudhan, CC BY-SA 4.0)

# K8s Service



- Cluster
- Services
- Pods
- Nodes

# K8s Service

---

- Service is a method for exposing a network application that is running as one or more pods in the cluster.
  - The Service API is an abstraction to expose groups of Pods over a network.
  - Run application code in Pods
- **apiVersion:** v1
  - **kind:** Service
  - **metadata:**
    - **name:** my-service
  - **spec:**
    - **selector:**
      - **app.kubernetes.io/name:** MyApp
    - **ports:**
      - - **protocol:** TCP
      - **port:** 80
      - **targetPort:** 9376

# Kubectl Commands

- kubectl get nodes

```
PS C:\Users\BITS\Desktop\Shwetha\2024-25\CC\FallSem-Oct2024\lab> kubectl get nodes
NAME             STATUS    ROLES    AGE   VERSION
docker-desktop   Ready     control-plane  19h   v1.30.2
```

- Status of node: Ready, NotReady, unknown
  - **Ready:** The node is running healthy and pods can be run inside.
  - **NotReady:** The node is not yet ready to run the pods
  - **Unknown:** If the node is not responding. If the master node cannot communicate with that node



# Kubectl Commands

---

- `kubectl create deployment hello-node --image=registry.k8s.io/e2e-test-images/agnhost:2.39 -- /agnhost netexec --http-port=8080`
- `kubectl get deployments`

```
PS C:\Users\BITS\Desktop\Shwetha\2024-25\CC\FallSem-Oct2024\lab> kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hello-node	1/1	1	1	17h

# Kubectl Commands

## kubectl get pods

```
PS C:\Users\BITS\Desktop\Shwetha\2024-25\CC\FallSem-Oct2024\lab> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-node-55fdcd95bf-22mh2	1/1	Running	0	18h

## kubectl get events

```
PS C:\Users\BITS\Desktop\Shwetha\2024-25\CC\FallSem-Oct2024\lab> kubectl get events
```

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
84s	Normal	Scheduled	pod/hello-node-55fdcd95bf-22mh2	Successfully assigned default/hello-node-55fdcd95bf-22mh2 to docker-desktop
83s	Normal	Pulling	pod/hello-node-55fdcd95bf-22mh2	Pulling image "registry.k8s.io/e2e-test-images/agnhost:2.39"
68s	Normal	Pulled	pod/hello-node-55fdcd95bf-22mh2	Successfully pulled image "registry.k8s.io/e2e-test-images/agnhost:2.39" in 14.666s (14.666s including wait ing). Image size: 51105200 bytes.
68s	Normal	Created	pod/hello-node-55fdcd95bf-22mh2	Created container agnhost
68s	Normal	Started	pod/hello-node-55fdcd95bf-22mh2	Started container agnhost
84s	Normal	SuccessfulCreate	replicaset/hello-node-55fdcd95bf	Created pod: hello-node-55fdcd95bf-22mh2
84s	Normal	ScalingReplicaSet	deployment/hello-node	Scaled up replica set hello-node-55fdcd95bf to 1

# Kubectl Commands

- Kubectl get logs <pod name>

```
PS C:\Users\BITS\Desktop\Shwetha\2024-25\CC\FallSem-Oct2024\lab> kubectl logs hello-node-55fdcd95bf-22mh2
1112 14:20:48.871286      1 log.go:195] Started HTTP server on port 8080
1112 14:20:48.872371      1 log.go:195] Started UDP server on port 8081
```

- Kubectl get services

```
PS C:\Users\BITS\Desktop\Shwetha\2024-25\CC\FallSem-Oct2024\lab> kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	6d23h
nginx-service-np	NodePort	10.108.167.246	<none>	8081:30107/TCP	83m

# YML Specification

---

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

# YML Specification

---

```
name: deploy a web server
k8s:
  api_version: v1
  namespace: my-namespace
  definition:
    kind: Deployment
    metadata:
      labels:
        app: nginx
    name: nginx-deploy
    spec:
      replicas: 1
      selector:
        matchLabels:
          app: nginx
    template:
      metadata:
        labels:
          app: nginx
```


# YML Specification

```
apiVersion: apps/v1 # Specifies the API version for the deployment resource.
kind: Deployment # Declares the type of Kubernetes object being created; in this case, a Deployment.
metadata:
  name: nginx-test-deployment # The name of the deployment. Used for identification within the cluster.
  labels: # Labels are key-value pairs for organizing and selecting resources.
    app: nginx # Label with key 'app' and value 'nginx'. Helps in identifying related resources.
spec:
  replicas: 1 # Specifies the number of desired pod replicas for this deployment.
  selector: # Defines how the Deployment finds which Pods to manage.
    matchLabels:
      app: nginx # The Deployment will manage Pods with this label.
  template: # Template section defines the Pod specification for the deployment.
    metadata:
      labels:
        app: nginx # Labels assigned to Pods created by this Deployment.
    spec: # The Pod specification, which defines the containers and other settings.
      containers: # List of containers to be run in the Pod.
        - name: nginx # The name of the container, useful for identifying the container within the Pod.
          image: nginx:latest # The container image to be used. 'nginx:latest' specifies the latest version of nginx.
          ports:
            - containerPort: 80 # The port that this container exposes. It maps to port 80 inside the container.
```

Note: This file is placed for understanding the terminology only. Does not have indentation correct. Hence, cannot be used with K8s  
apply/create to create the deployment


# K8s Service Types – Cluster IP

- Default Service Type
- Internal Communication – Inside Cluster
- Exposes an internal IP within the cluster
- No external access allowed




# K8s Service Types – NodePort

- Exposes a specific application to the outside world with a port on the node
- External access to the service
- Forwards the traffic from the specified port on each node to the corresponding port on the pods targeted by the service.
- When application are to be accessed from outside the cluster without requiring a load balancer
- Note: One service per port
- Node port is static – 30000 to 32767





# K8s Service Types – Load Balancer



- Provides external communication
- Exposes the application to the outside world – public IP
- Provides load balancing functions too
- E.g.: Web server

# Use cases of K8s

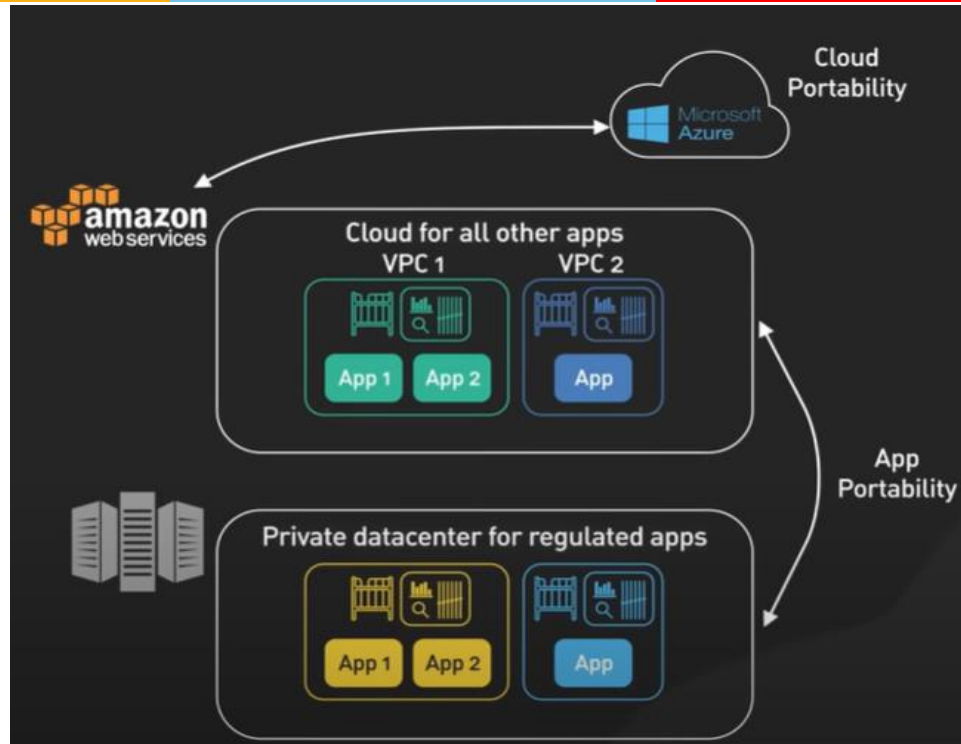


Image Courtesy: Bytebytego

- **Deployment and orchestration of microservices**
  - Ensures seamless deployment and communication.

# Use cases of K8s

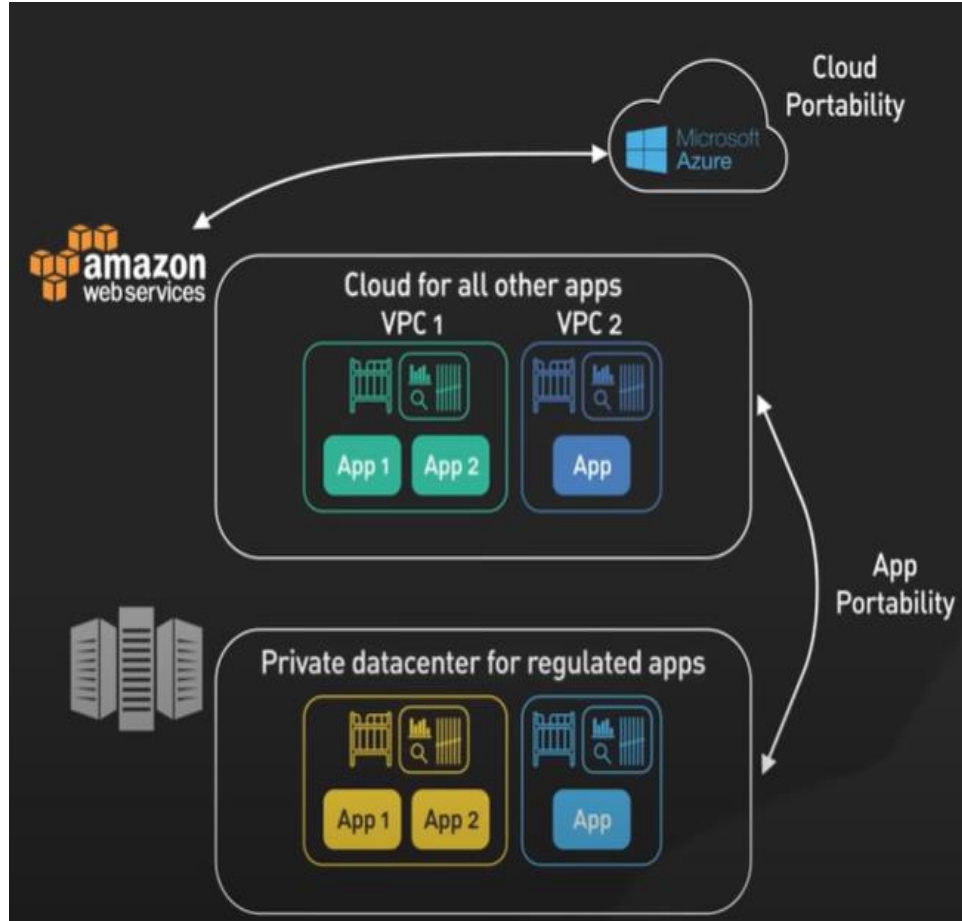


Image Courtesy: Bytebytego

- **Resource Optimization**
  - Dynamic allocation of resources based on workload
  - Maximizing system utilization
  - Reduces operational costs.
- **CI/CD Pipelines**
  - By automating deployment processes, K8s aids continuous integration and delivery.

# What K8s is not

---

- Does not deploy source code and does not build your application
- Does not provide application-level services
- Does not dictate logging, monitoring, or alerting solutions

# Summary

---

- Need for Container Orchestration
- K8s Orchestration Features & Functionalities
- K8s Objects – Nodes, Cluster, Pods, Deployments, Services
- K8s Service Types – ClusterIP, NodePort, LoadBalancer

Refer: [Overview | Kubernetes](#)

# Lab – nginx Service on K8s Cluster

---

Deploy nginx web service on <https://labs.play-with-k8s.com/>

1. Create K8s cluster with two nodes
2. Experiment with deployment and different service types with two nodes on a K8s.
3. Capture all your observations and discuss with peers and instructor.

nginx deployment – nginx.yml

Nginx service as ClusterIP – clusteripservice.yml

Nginx service as NodePort – nodeport.yml

Nginx service as Loadbalancer – lb.yml

# Lab – K8s Cluster creation with Two Nodes

Get into <https://labs.play-with-k8s.com/>

Create two instances here

On the first instance – (master node)

1. Initializes cluster master node:

```
kubeadm init --apiserver-advertise-address $(hostname -i) --pod-network-cidr 10.5.0.0/16
```

2. Initialize cluster networking:

```
kubectl apply -f https://raw.githubusercontent.com/cloudnativelabs/kube-router/master/daemonset/kubeadm-kuberouter.yaml
```

3. Create an nginx deployment:

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/website/master/content/en/examples/application/nginx-app.yaml
```

On the second instance terminal (worker node)

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 192.168.0.18:6443 --token 4goo7z.k6laq8hufybdrnen \
--discovery-token-ca-cert-hash
sha256:2c4b669dcab8e1342befb6a22ff882dea58a501f5a8b13d5aa260f7f5a3298fd
```

# nginx Deployment

---

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.23.3 #latest version on DockerHub
          ports:
            - containerPort: 80
```

kubectl apply -f nginx.yml

kubectl get deployments

kubectl get pods

kubectl logs <pod name>



# nginx clusterIP Service

---

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8081 -> Service accessible on 8081
      targetPort: 80 -> Container inside the pod uses
port 80
```

Kubectl apply -f  
clusteripservice.yml

Kubectl get svc

Note the cluster IP from the output

Go to other instance and execute  
curl <cluster ip of the service>:  
8081

# nginx NodePort Service

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service-np
spec:
  type: NodePort
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8081 Service accessible on 8081
      targetPort: 80 -> Container inside the pod uses
port 80
      nodePort: 32002 ->Static port assigned to the
node
```

Kubectl apply -f nodeport.yml

Kubectl get svc

Note the cluster IP from the output

Go to other instance and give  
curl <cluster ip of the service>:  
8081

Also  
Curl <node ip>:32002

# nginx LoadBalancer Service

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service-lb
spec:
  type: LoadBalancer
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8081
      targetPort: 80
      nodePort: 32002
```

Kubectrl apply -f lb.yml

Kubectrl get svc

Note the cluster IP from the output

Go to other instance and give  
curl <cluster ip of the service>: 8081

Also

Curl <node ip>:32002

What additional things can you do on this service ?

---

# **Additional Slides**

# Additional Features of K8s

---

- **Configuration and Secret management**
  - Store and manage sensitive information, such as passwords, OAuth tokens, and SSH keys.
  - Deploy and update secrets and application configuration without rebuilding your container images, and without exposing secrets in your stack configuration.
- **Batch Execution**
  - Manage batch and CI workloads

# K8s Architecture – Control Plane

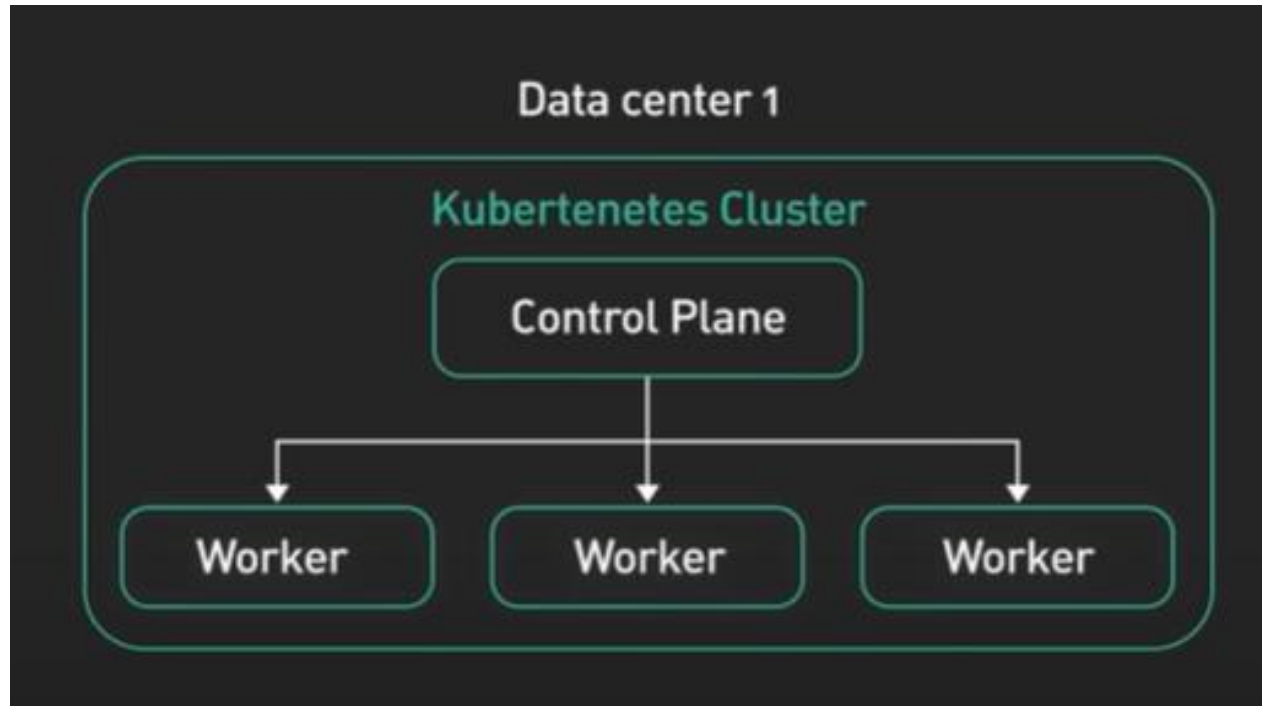


Image Courtesy: Bytebytego

## Two Core Parts

### 1. Control Plane (Master)


Responsible for managing the state of cluster

### 2. Worker Node(s)

Run the containerized applications workloads

# K8s Architecture – Control Plane

- **API Server:**
  - Primary interface between the control plane and rest of the cluster.
- **Etcd:**
  - Stores the cluster's persistent state
  - What resources are available
  - Health information of cluster
  - Other components use /update this information about the cluster



# K8s Architecture – Control Plane

- **Scheduler:**
  - Responsible for scheduling pods on the worker nodes
  - Uses available and required resources information for the pods to perform its job.

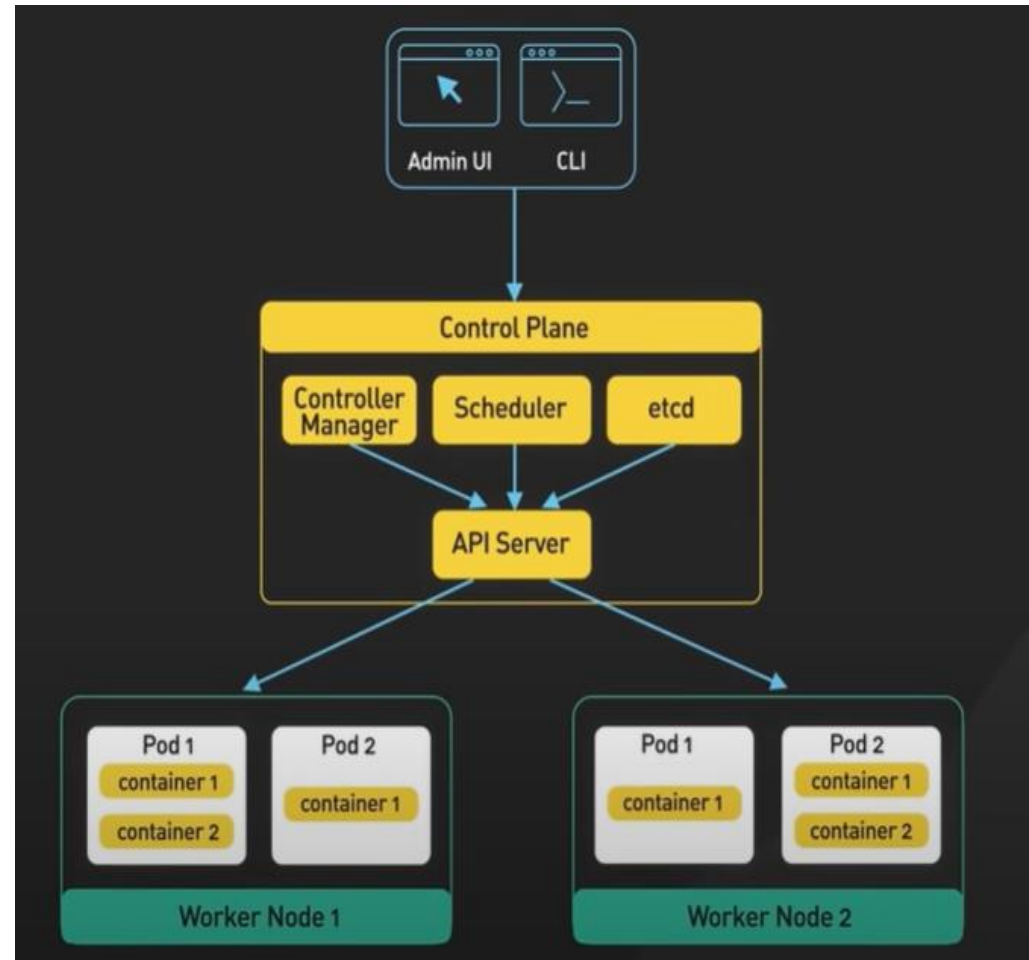


Image Courtesy: Bytebytego



# K8s Architecture – Control Plane

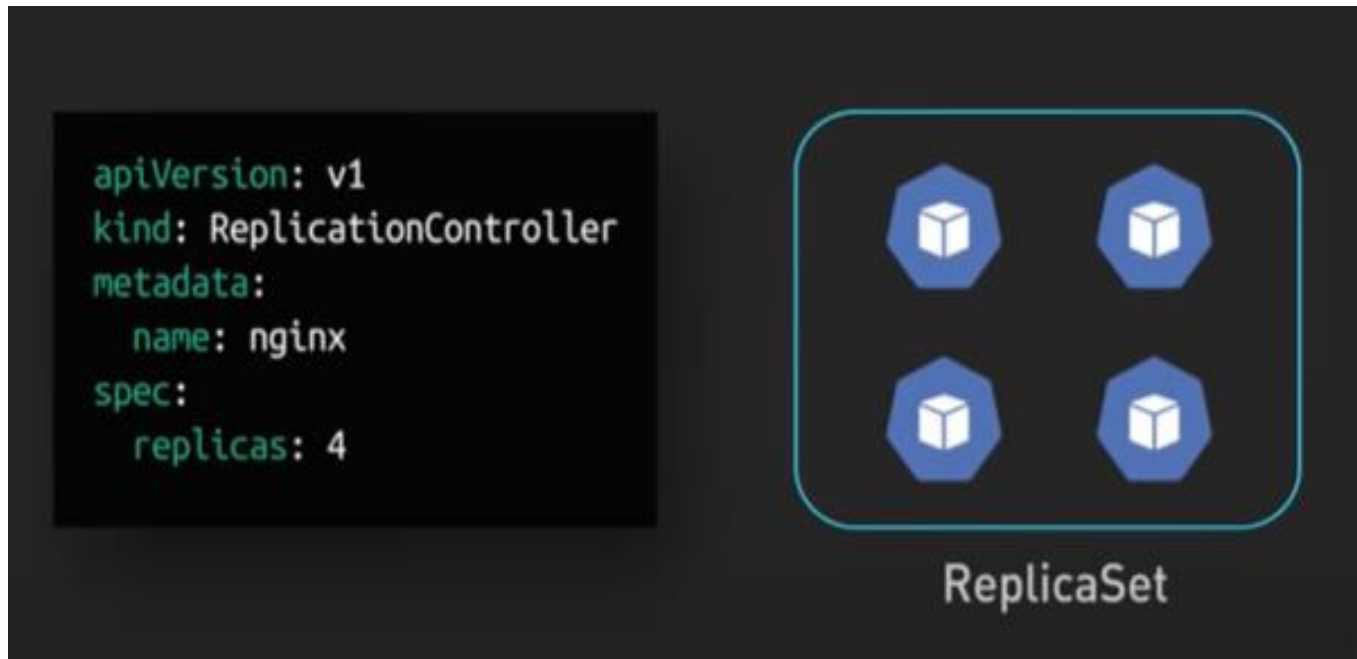


Image Courtesy: Bytebytego

- **Controller Manager**
  - Responsible for running those containers that manage the state of the cluster.
  - E.g: ReplicationController

# K8s Architecture – Worker Node

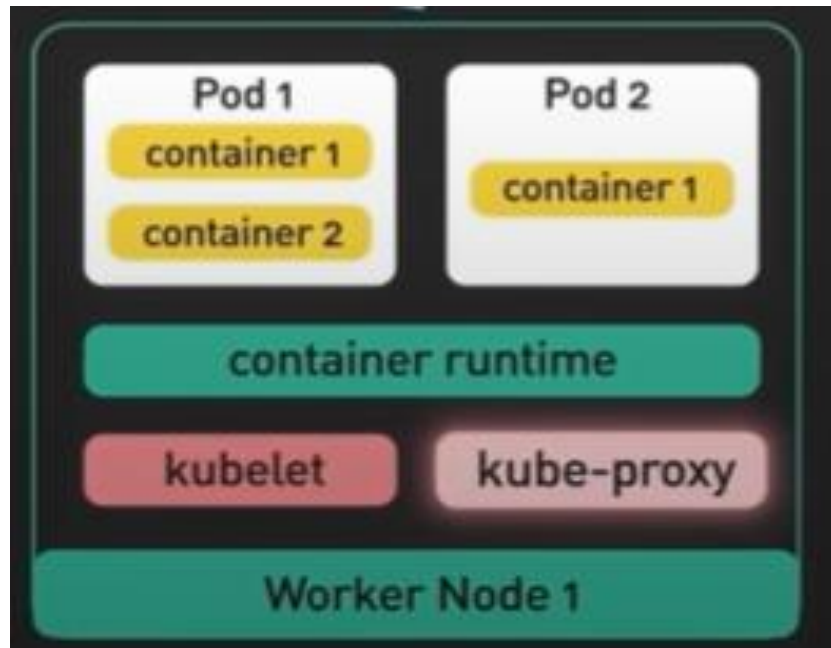


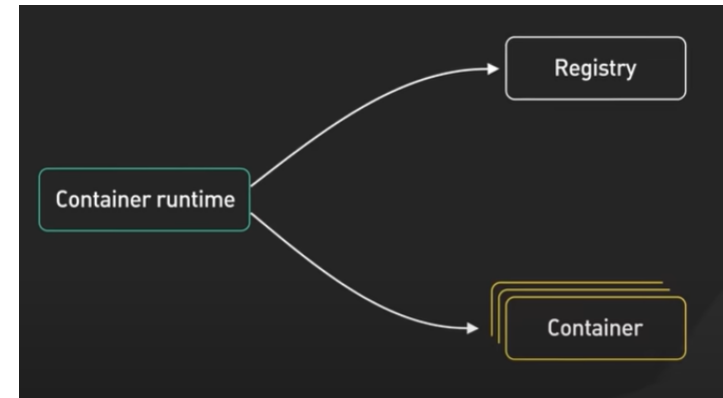
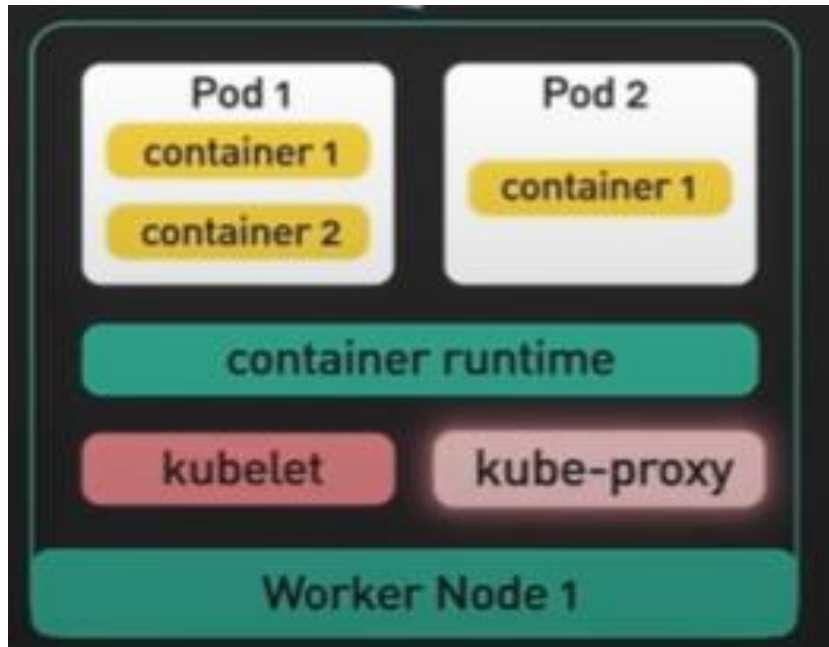
Image Courtesy: Bytebytego

## Kubelet

- Daemon running on each worker node
- Communicates with control plane.
- Receives information about which pods to run on the node and desired state of the pod is maintained.

**KubeProxy:** Routes the traffic to the correct pod, load balancing for even distribution of traffic across the pods

# K8s Architecture – Worker Node



All Image Courtesy: Bytebytego.com

## Container Runtime

Runs the container on the worker nodes